

In [69]:

```
import numpy, pandas, scipy
#Rescaling the data
from sklearn.preprocessing import MinMaxScaler
df=pandas.read_csv("C:/Users/ABHISHEK/Desktop/GITAM ML/GITAM ML LAB/redwinequality.csv")
df.shape
```

Out[69]:

(1599, 12)

In [45]:

df.head()

Out[45]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9

In [70]:

df.describe()

Out[70]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000

In [73]:

```
#Seperating data into input and output components
x=df.iloc[:,0:11]
y=df.iloc[:,11:]
df.describe()
```

Out[73]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
<b>count</b>	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
<b>mean</b>	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467000
<b>std</b>	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895000
<b>min</b>	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
<b>25%</b>	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000
<b>50%</b>	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000
<b>75%</b>	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000
<b>max</b>	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000

In [74]:

```
x.head(2)
```

Out[74]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
<b>0</b>	7.4	0.70	0.0	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9
<b>1</b>	7.8	0.88	0.0	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9

In [80]:

```
#if no argument then prints all y
y.head(3)
```

Out[80]:

	quality
<b>0</b>	5
<b>1</b>	5
<b>2</b>	5

In [81]:

```
#defining the scaler
scaler=MinMaxScaler(feature_range=(0,1))
print(scaler)
print(type(scaler))
```

```
MinMaxScaler(copy=True, feature_range=(0, 1))
<class 'sklearn.preprocessing.data.MinMaxScaler'>
```

In [82]:

```
rescaledX=scaler.fit_transform(x)
numpy.set_printoptions(precision=2) #Setting precision for the output
print(rescaledX[0:5,:])
```

```
[[0.25 0.4  0.    0.07 0.11 0.14 0.1  0.57 0.61 0.14 0.15]
 [0.28 0.52 0.    0.12 0.14 0.34 0.22 0.49 0.36 0.21 0.22]
 [0.28 0.44 0.04 0.1  0.13 0.2  0.17 0.51 0.41 0.19 0.22]
 [0.58 0.11 0.56 0.07 0.11 0.23 0.19 0.58 0.33 0.15 0.22]
 [0.25 0.4  0.    0.07 0.11 0.14 0.1  0.57 0.61 0.14 0.15]]
```

In [83]:

```
rescaledX.max()
```

Out[83]:

```
1.0000000000000002
```

In [84]:

```
from sklearn.preprocessing import StandardScaler
#code for standardizing the data
scaler=StandardScaler().fit(x)
print(scaler)
print(type(scaler))
```

```
StandardScaler(copy=True, with_mean=True, with_std=True)
<class 'sklearn.preprocessing.data.StandardScaler'>
```

In [85]:

```
rescaledX=scaler.transform(x)
print(rescaledX[0:5,:])
```

```
[[-0.53  0.96 -1.39 -0.45 -0.24 -0.47 -0.38  0.56  1.29 -0.58 -0.96]
 [-0.3   1.97 -1.39  0.04  0.22  0.87  0.62  0.03 -0.72  0.13 -0.58]
 [-0.3   1.3  -1.19 -0.17  0.1  -0.08  0.23  0.13 -0.33 -0.05 -0.58]
 [ 1.65 -1.38  1.48 -0.45 -0.26  0.11  0.41  0.66 -0.98 -0.46 -0.58]
 [-0.53  0.96 -1.39 -0.45 -0.24 -0.47 -0.38  0.56  1.29 -0.58 -0.96]]
```

In [86]:

```
rescaledX.mean()
```

Out[86]:

2.5450106517076758e-15

In [87]:

```
from sklearn.preprocessing import Normalizer
#code for normalizing the data
scaler=Normalizer().fit(x)
normalizedX=scaler.transform(x)
print(normalizedX[0:5,:])
```

```
[[1.95e-01 1.85e-02 0.00e+00 5.01e-02 2.00e-03 2.90e-01 8.97e-01 2.63e-02
 9.26e-02 1.48e-02 2.48e-01]
 [1.07e-01 1.21e-02 0.00e+00 3.57e-02 1.35e-03 3.44e-01 9.21e-01 1.37e-02
 4.40e-02 9.35e-03 1.35e-01]
 [1.35e-01 1.32e-02 6.95e-04 3.99e-02 1.60e-03 2.60e-01 9.38e-01 1.73e-02
 5.66e-02 1.13e-02 1.70e-01]
 [1.74e-01 4.36e-03 8.72e-03 2.96e-02 1.17e-03 2.65e-01 9.34e-01 1.55e-02
 4.92e-02 9.03e-03 1.53e-01]
 [1.95e-01 1.85e-02 0.00e+00 5.01e-02 2.00e-03 2.90e-01 8.97e-01 2.63e-02
 9.26e-02 1.48e-02 2.48e-01]]
```

In [88]:

```
row1 = [1.95152519e-01, 1.84603734e-02, 0.00000000e+00, 5.01067279e-02,
2.00426911e-03, 2.90091582e-01, 8.96646709e-01, 2.63139437e-02,
 9.25655867e-02, 1.47682987e-02, 2.47896443e-01]
res = sum(map(lambda i : i * i, row1))
res
```

Out[88]:

0.99999999996141229

In [91]:

```
from sklearn.preprocessing import Binarizer
#code for binarizing the data
binarizer=Binarizer(threshold=0.5).fit(x)
binaryX=binarizer.transform(x)
print(binaryX[0:5,:])
```

```
[[1. 1. 0. 1. 0. 1. 1. 1. 1. 1. 1.]
 [1. 1. 0. 1. 0. 1. 1. 1. 1. 1. 1.]
 [1. 1. 0. 1. 0. 1. 1. 1. 1. 1. 1.]
 [1. 0. 1. 1. 0. 1. 1. 1. 1. 1. 1.]
 [1. 1. 0. 1. 0. 1. 1. 1. 1. 1. 1.]]
```

In [92]:

```
print(binaryX.min())  
print(binaryX.max())
```

```
0.0  
1.0
```

In [94]:

```
print(x.mean(axis=0))
```

```
fixed acidity      8.319637  
volatile acidity   0.527821  
citric acid        0.270976  
residual sugar     2.538806  
chlorides          0.087467  
free sulfur dioxide 15.874922  
total sulfur dioxide 46.467792  
density            0.996747  
pH                 3.311113  
sulphates          0.658149  
alcohol            10.422983  
dtype: float64
```

In [97]:

```
print(x.std(axis=0))
```

```
fixed acidity      1.741096  
volatile acidity   0.179060  
citric acid        0.194801  
residual sugar     1.409928  
chlorides          0.047065  
free sulfur dioxide 10.460157  
total sulfur dioxide 32.895324  
density            0.001887  
pH                 0.154386  
sulphates          0.169507  
alcohol            1.065668  
dtype: float64
```

In [98]:

```
from sklearn.preprocessing import scale  
#code for mean removal  
data_standardized=scale(x)
```

In [100]:

```
print(data_standardized.mean(axis=0))  
print(data_standardized.std(axis=0))
```

```
[ 3.55e-16  1.73e-16 -8.89e-17 -1.24e-16  3.91e-16 -6.22e-17  4.44e-17  
 2.36e-14  2.86e-15  6.75e-16  7.11e-17]  
[1.  1.  1.  1.  1.  1.  1.  1.  1.  1.]
```

In [101]:

```
# Method1: One-hot encoding directly from pandas using "pandas.get_dummies"
ids = [1, 2, 3, 4, 5, 6, 7]
place = ['Bangladesh', 'France', 'Germany', 'India', 'Nepal', 'Spain', 'Srilanka']
df = pandas.DataFrame(list(zip(ids, place)),
    columns=['Ids', 'Place'])
print(df)
print()
y = pandas.get_dummies(df.Place, prefix='P')
print(y)
```

	Ids	Place
0	1	Bangladesh
1	2	France
2	3	Germany
3	4	India
4	5	Nepal
5	6	Spain
6	7	Srilanka

	P_Bangladesh	P_France	P_Germany	P_India	P_Nepal	P_Spain	P_Srilanka
0	1	0	0	0	0	0	0
1	0	1	0	0	0	0	0
2	0	0	1	0	0	0	0
3	0	0	0	1	0	0	0
4	0	0	0	0	1	0	0
5	0	0	0	0	0	1	0
6	0	0	0	0	0	0	1

In [108]:

```
# Method2: One-hot encoding using "LabelBinarizer".
from sklearn.preprocessing import LabelBinarizer
ids = [11, 22, 33, 44, 55, 66, 77]
countries = ['France', 'Spain', 'Germany', 'France', 'Spain']
df = pandas.DataFrame(list(zip(ids, countries)),
    columns=['Ids', 'Countries'])
y = LabelBinarizer().fit_transform(df.Countries)
print(y)
```

```
[[1 0 0]
 [0 0 1]
 [0 1 0]
 [1 0 0]
 [0 0 1]]
```

In [110]:

```
# Method3: One-hot encoding using "OneHotEncoder".
from sklearn.preprocessing import OneHotEncoder
x = [[11, "France"], [22, "Spain"], [33, "Germany"], [44, "France"], [55, "Spain"]]
# x = [[1, 'Bangladesh'], [2, 'France'], [3, 'Germany']]
y = OneHotEncoder().fit_transform(x).toarray()
print(y)
```

```
[[1. 0. 0. 0. 0. 1. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 1.]
 [0. 0. 1. 0. 0. 0. 1. 0.]
 [0. 0. 0. 1. 0. 1. 0. 0.]
 [0. 0. 0. 0. 1. 0. 0. 1.]]
```

In [116]:

```
#Label Encoding
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
input_classes = ['suzuki', 'ford', 'suzuki', 'toyota', 'ford', 'bmw']
label_encoder.fit(input_classes)
print("\nClass mapping: ")
for i, item in enumerate(label_encoder.classes_):
    print(item, '-->', i)
```

Class mapping:

```
bmw --> 0
ford --> 1
suzuki --> 2
toyota --> 3
```

In [119]:

```
# Encoded Lables
labels = ['toyota', 'ford', 'bmw', 'suzuki', 'bmw']
encoded_labels = label_encoder.transform(labels)
print("\nLabels =", labels)
print("Encoded labels =", list(encoded_labels))
```

```
Labels = ['toyota', 'ford', 'bmw', 'suzuki', 'bmw']
Encoded labels = [3, 1, 0, 2, 0]
```

In [122]:

```
# Decoding the Lables
encoded_labels = [3, 2, 0, 2, 1]
decoded_labels = label_encoder.inverse_transform(encoded_labels)
print("\nEncoded labels =", encoded_labels)
print("Decoded labels =", list(decoded_labels))
```

```
Encoded labels = [3, 2, 0, 2, 1]
Decoded labels = ['toyota', 'suzuki', 'bmw', 'suzuki', 'ford']
```

In [ ]: